

Reliable model evaluation

The Transformer uses attention to combine information.

The test used 24 GB of GPU memory and 32 examples per batch.

Reported accuracy was 92.5 percent on the held-out dataset.

Objective: $L = \sum(x_i * y_i)$

The figure below is part of the original PDF.

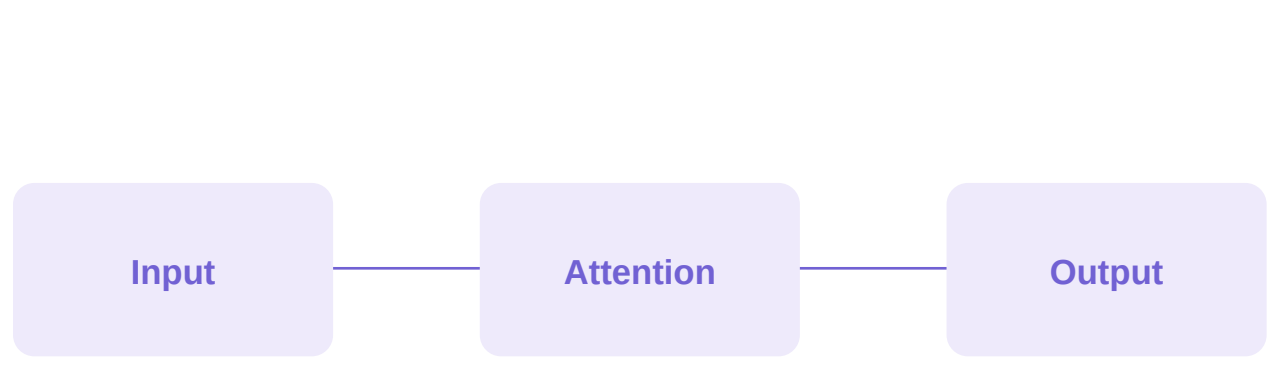


Figure 1. A simple attention pipeline. Preserve the visual unchanged.

The article does not report GPU model, training time, or dataset license.